

loan-approval-prediction

October 22, 2024

1 Loan Approval Prediction Using Machine Learning

2 Importing necessary libraries

```
[1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, \
    classification_report
import matplotlib.pyplot as plt
import seaborn as sns
from imblearn.over_sampling import SMOTE
from xgboost import XGBClassifier
from sklearn.impute import SimpleImputer
```

3 Loading the dataset

```
[2]: df = pd.read_csv('C:/Users/ASUS/Desktop/Power BI Practice/loan_data.csv')
df.head(10)
```

```
[2]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0.0	Graduate	No	
1	LP001003	Male	Yes	1.0	Graduate	No	
2	LP001005	Male	Yes	0.0	Graduate	Yes	
3	LP001006	Male	Yes	0.0	Not Graduate	No	
4	LP001008	Male	No	0.0	Graduate	No	
5	LP001011	Male	Yes	2.0	Graduate	Yes	
6	LP001013	Male	Yes	0.0	Not Graduate	No	
7	LP001014	Male	Yes	3.0	Graduate	No	
8	LP001018	Male	Yes	2.0	Graduate	No	
9	LP001020	Male	Yes	1.0	Graduate	No	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	NaN	360.0	

1	4583	1508.0	128.0	360.0
2	3000	0.0	66.0	360.0
3	2583	2358.0	120.0	360.0
4	6000	0.0	141.0	360.0
5	5417	4196.0	267.0	360.0
6	2333	1516.0	95.0	360.0
7	3036	2504.0	158.0	360.0
8	4006	1526.0	168.0	360.0
9	12841	10968.0	349.0	360.0

	Credit_History	Property_Area	Loan_Status
0	1.0	Urban	Y
1	1.0	Rural	N
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y
5	1.0	Urban	Y
6	1.0	Urban	Y
7	0.0	Semiurban	N
8	1.0	Urban	Y
9	1.0	Semiurban	N

4 Basic data exploration

```
[3]: print(df.info())    # Display the data types and missing values
      print(df.describe()) # Show basic statistics of numerical columns
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 598 entries, 0 to 597
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID               598 non-null   object
1   Gender                598 non-null   object
2   Married               598 non-null   object
3   Dependents            586 non-null   float64
4   Education              598 non-null   object
5   Self_Employed         598 non-null   object
6   ApplicantIncome       598 non-null   int64
7   CoapplicantIncome     598 non-null   float64
8   LoanAmount            577 non-null   float64
9   Loan_Amount_Term      584 non-null   float64
10  Credit_History         549 non-null   float64
11  Property_Area          598 non-null   object
12  Loan_Status            598 non-null   object
dtypes: float64(5), int64(1), object(7)
```

memory usage: 60.9+ KB

None

	Dependents	ApplicantIncome	CoapplicantIncome	LoanAmount	\
count	586.000000	598.000000	598.000000	577.000000	
mean	0.755973	5292.252508	1631.499866	144.968804	
std	1.007751	5807.265364	2953.315785	82.704182	
min	0.000000	150.000000	0.000000	9.000000	
25%	0.000000	2877.500000	0.000000	100.000000	
50%	0.000000	3806.000000	1211.500000	127.000000	
75%	1.750000	5746.000000	2324.000000	167.000000	
max	3.000000	81000.000000	41667.000000	650.000000	

	Loan_Amount_Term	Credit_History
count	584.000000	549.000000
mean	341.917808	0.843352
std	65.205994	0.363800
min	12.000000	0.000000
25%	360.000000	1.000000
50%	360.000000	1.000000
75%	360.000000	1.000000
max	480.000000	1.000000

5 Handle missing values (if any)

```
[4]: # Fill missing values with median or mode depending on the column type
df['LoanAmount'] = df['LoanAmount'].fillna(df['LoanAmount'].median())
df['Loan_Amount_Term'] = df['Loan_Amount_Term'].fillna(df['Loan_Amount_Term'].
    ↪median())
df['Credit_History'] = df['Credit_History'].fillna(df['Credit_History'].
    ↪mode()[0])

df['Gender'] = df['Gender'].fillna(df['Gender'].mode()[0])
df['Married'] = df['Married'].fillna(df['Married'].mode()[0])
df['Dependents'] = df['Dependents'].fillna(df['Dependents'].mode()[0])
df['Self_Employed'] = df['Self_Employed'].fillna(df['Self_Employed'].mode()[0])
```

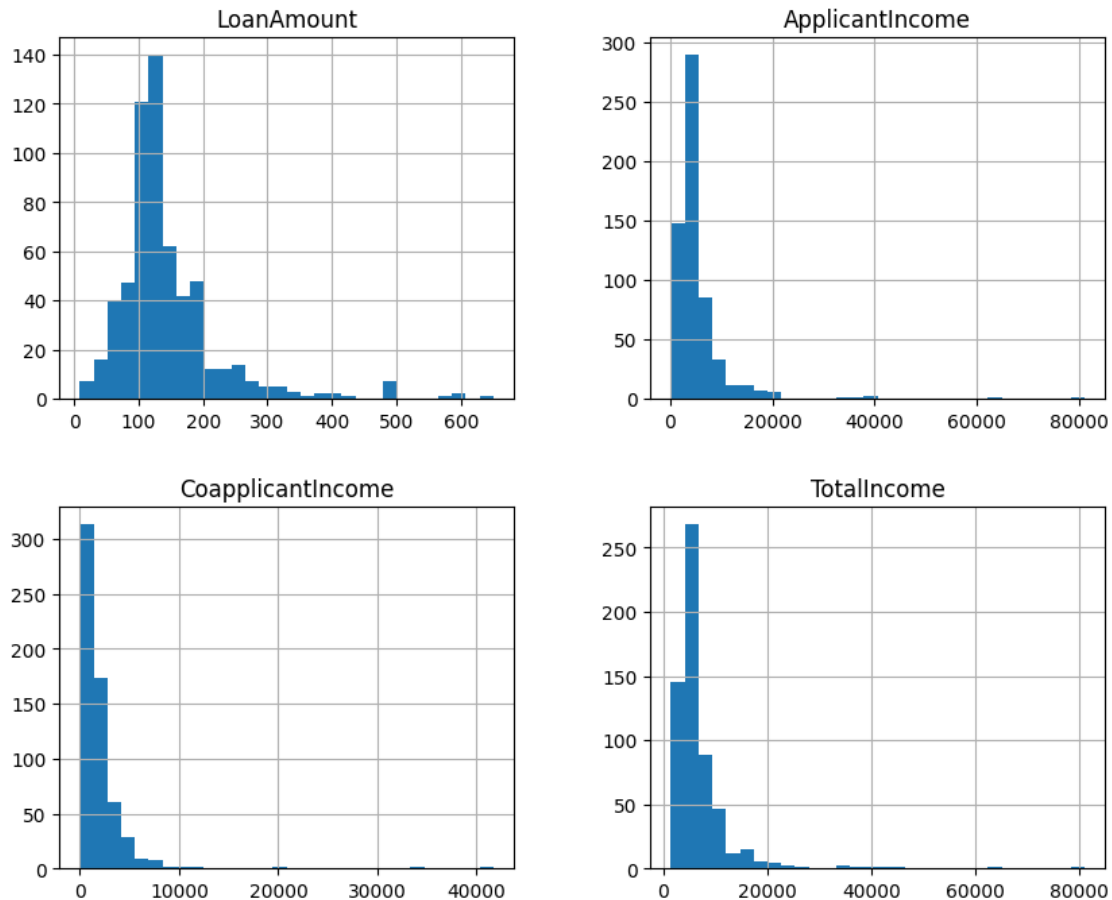
6 Feature engineering - Adding log-transformed values for skewed numerical columns

```
[5]: # Log-transform skewed numerical columns to normalize them
df['LoanAmount_log'] = np.log(df['LoanAmount'] + 1) # Adding 1 to avoid log(0)
df['TotalIncome'] = df['ApplicantIncome'] + df['CoapplicantIncome']
df['TotalIncome_log'] = np.log(df['TotalIncome'] + 1)
```

7 Visualizing distributions of numerical features

```
[6]: df[['LoanAmount', 'ApplicantIncome', 'CoapplicantIncome', 'TotalIncome']].  
      hist(bins=30, figsize=(10, 8))  
plt.suptitle('Distributions of Loan Amount and Income')  
plt.show()
```

Distributions of Loan Amount and Income




8 Bar plot of Loan_Status counts to visualize class distribution

```
[7]: plt.figure(figsize=(6, 4))  
sns.countplot(x='Loan_Status', data=df)  
plt.title('Loan Approval Distribution')  
plt.show()
```



9 Correlation heatmap - only for numerical columns

```
[8]: # Correlation analysis of numerical features to identify relationships
plt.figure(figsize=(10, 8))
numerical_df = df.select_dtypes(include=[np.number]) # Select only numerical
↳ columns
corr_matrix = numerical_df.corr()
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix')
plt.show()
```



10 Encode categorical variables

```
[9]: # Label encode categorical columns
encoder = LabelEncoder()
df['Gender'] = encoder.fit_transform(df['Gender'])
df['Married'] = encoder.fit_transform(df['Married'])
df['Education'] = encoder.fit_transform(df['Education'])
df['Self_Employed'] = encoder.fit_transform(df['Self_Employed'])
df['Property_Area'] = encoder.fit_transform(df['Property_Area'])
df['Loan_Status'] = encoder.fit_transform(df['Loan_Status']) # Target variable

# Replace '3+' in Dependents with 3 and convert to integer type
df['Dependents'] = df['Dependents'].replace('3+', 3).astype(int)
```

11 Prepare the data for modeling

```
[10]: # Select features and target variable for modeling
X = df[['Gender', 'Married', 'Dependents', 'Education', 'Self_Employed',
        ↪ 'LoanAmount_log',
        'Loan_Amount_Term', 'Credit_History', 'Property_Area',
        ↪ 'TotalIncome_log']]
y = df['Loan_Status']
```

12 Split the data into training and testing sets

```
[11]: # Use train_test_split to create training and testing datasets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
        ↪ random_state=42)
```

13 Model training - RandomForestClassifier

```
[12]: # Train a Random Forest classifier on the training data
model = RandomForestClassifier(random_state=42)
model.fit(X_train, y_train)
```

```
[12]: RandomForestClassifier(random_state=42)
```

14 Make predictions and evaluate the model

```
[13]: y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy * 100:.2f}%') # Printing accuracy of the model
```

Accuracy: 77.22%

15 Display the confusion matrix and classification report

```
[14]: conf_matrix = confusion_matrix(y_test, y_pred)
print('Confusion Matrix:\n', conf_matrix)
print('Classification Report:\n', classification_report(y_test, y_pred))
```

Confusion Matrix:

```
[[ 25  31]
 [ 10 114]]
```


Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.71	0.45	0.55	56
1	0.79	0.92	0.85	124
accuracy			0.77	180
macro avg	0.75	0.68	0.70	180
weighted avg	0.76	0.77	0.75	180

16 Visualize the confusion matrix

```
[15]: plt.figure(figsize=(4, 2))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



17 Handle class imbalance using SMOTE

```
[16]: smote = SMOTE(random_state=42)
X_res, y_res = smote.fit_resample(X, y)
```

18 Train an XGBoost model for comparison

```
[17]: model_xgb = XGBClassifier(random_state=42) # Removed use_label_encoder
model_xgb.fit(X_res, y_res)
y_pred_xgb = model_xgb.predict(X_test)

xgb_accuracy = accuracy_score(y_test, y_pred_xgb)
```



```
print(f'XGBoost Accuracy: {xgb_accuracy * 100:.2f}%')
```

XGBoost Accuracy: 100.00%

```
[18]: # Display confusion matrix and classification report for XGBoost
conf_matrix_xgb = confusion_matrix(y_test, y_pred_xgb)
print('Confusion Matrix (XGBoost):\n', conf_matrix_xgb)
print('Classification Report (XGBoost):\n', classification_report(y_test,
↪y_pred_xgb))
```

Confusion Matrix (XGBoost):

```
[[ 56   0]
```


```
 [  0 124]]
```

Classification Report (XGBoost):

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56
1	1.00	1.00	1.00	124
accuracy			1.00	180
macro avg	1.00	1.00	1.00	180
weighted avg	1.00	1.00	1.00	180

19 Visualize the confusion matrix for XGBoost

```
[19]: plt.figure(figsize=(4, 2))
sns.heatmap(conf_matrix_xgb, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.title('Confusion Matrix (XGBoost)')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



20 Feature Importance

```
[20]: # Visualize the importance of different features used in the Random Forest model
feature_importances = model.feature_importances_
features = X.columns
importance_df = pd.DataFrame({'Feature': features, 'Importance':
    ↪feature_importances})
importance_df = importance_df.sort_values(by='Importance', ascending=False)
plt.figure(figsize=(10, 6))
sns.barplot(x='Importance', y='Feature', data=importance_df, hue='Feature',
    ↪palette='coolwarm', dodge=False, legend=False)
plt.title('Feature Importances')
plt.show()
```

